Advanced Machine Learning Architectures for Exoplanetary Digital Twins: Strategic Blueprints for Astrobiological Climatology

Introduction to Predictive Exoclimatology and Digital Twins

The transition from purely observational astronomy to predictive exoclimatology represents one of the most profound paradigm shifts in modern astrophysics. Historically, the search for exoplanets was dominated by the statistical accumulation of discoveries, seeking to answer whether planetary bodies beyond our solar system existed in significant numbers. With missions such as Kepler, the Transiting Exoplanet Survey Satellite (TESS), and the James Webb Space Telescope (JWST) operating at unprecedented resolutions, the discipline has evolved into the intricate characterization of these distant worlds.[1] Consequently, astrobiology now faces a formidable data challenge: identifying which of the thousands of confirmed exoplanets possesses the necessary physical, atmospheric, and climatic conditions to harbor life.[4]

To address the sheer volume and dimensionality of this data, researchers increasingly rely on the concept of the "digital twin." A digital twin, in the context of planetary science, is a highly sophisticated, continuously refined virtual replica of a physical system. It allows scientists to run simulations and alter boundary conditions to observe how an exoplanet's climate, ocean heat transport, and atmospheric chemistry might evolve over geological timescales.[4] For exoplanet habitability, constructing these digital twins requires the synthesis of orbital mechanics, stellar irradiation models, and fluid dynamics, traditionally executed through computationally expensive General Circulation Models (GCMs).[7]

However, running a single planetary configuration through a classical three-dimensional GCM can require thousands of compute hours, rendering large-scale parameter sweeps impossible.[10] The integration of Artificial Intelligence (AI) and Machine Learning (ML) architectures provides a profound solution to this computational bottleneck. By replacing or augmenting traditional physics-based solvers with machine learning emulators—such as Extreme Learning Machines, Physics-Informed Neural Networks, and Generative Adversarial Networks—researchers can achieve predictions of planetary temperature distributions, atmospheric spectra, and habitability indices at a fraction of the computational cost.[2]

The following analysis provides an exhaustive exploration of how state-of-the-art machine learning paradigms can be applied to the construction of digital twins for exoplanet habitability.

Tailored specifically for advanced computational initiatives, such as Category 3 of the HACK-4-SAGES astrobiology hackathon, this report delineates the theoretical astrophysics underpinning planetary climates, evaluates the most effective neural architectures for emulation, dissects the structure of available astronomical datasets, and proposes rigorous, implementable project blueprints.[4] For students with advanced backgrounds in mathematics and computer science—such as those at the Warsaw University of Technology's Faculty of Mathematics and Information Science—this intersection of differential equations, statistics, and deep learning represents the ideal domain for rapid, high-impact algorithmic development. The synthesis of these domains demonstrates how data-driven architectures can fundamentally accelerate our understanding of habitability across the cosmos.

Theoretical Astrophysics and Climate Dynamics

To effectively train machine learning models, one must first establish the underlying physical mechanics that dictate exoplanetary climates. The boundary conditions for any planetary digital twin are established by stellar parameters, orbital mechanics, and the thermodynamic properties of the planet's surface and atmosphere.[2] A machine learning model is only as robust as the mathematical physics it attempts to emulate; thus, a deep understanding of these planetary equations is a prerequisite for feature engineering.

The Habitable Zone and Equilibrium Temperature

The fundamental metric of habitability is the theoretical capacity of a planet to maintain liquid water on its surface. This is colloquially referred to as the Habitable Zone (HZ). While classical definitions of the HZ are overly simplistic, they remain a necessary baseline for filtering vast astronomical datasets.[15] The primary indicator of a planet's surface conditions is its equilibrium temperature, assuming a uniform redistribution of heat. This temperature, T_{eq}, is derived from the Stefan-Boltzmann law and is defined mathematically as:

$$T_{eq} = \left(\frac{S(1-A)}{4\sigma} \right)^{1/4}$$

where S represents the incident stellar flux, A represents the planetary Bond albedo (the fraction of electromagnetic energy reflected back into space), and σ is the Stefan-Boltzmann constant.[4] While this equation provides an idealized baseline, actual surface temperatures are heavily modulated by the greenhouse effect, atmospheric density, and circulation patterns.[18]

For stars of different spectral types, particularly M-dwarfs which are highly abundant and frequently host terrestrial planets, the boundaries of the HZ are highly sensitive to stellar effective temperature. Advanced empirical models utilize polynomial equations to define these boundaries precisely. The inner and outer edges of the habitable zone can be calculated by

scaling the solar constant using specific coefficients derived from 3D climate simulations.[19] These coefficients account for phenomena such as the runaway greenhouse effect at the inner edge and the maximum carbon dioxide greenhouse limit at the outer edge.[21] Machine learning models tasked with identifying habitable candidates must encode these non-linear boundary conditions to accurately differentiate between a temperate rocky world and a barren, irradiated core.[22]

Orbital Mechanics and Tidal Locking Dynamics

A critical consideration for modern exoclimatology is the phenomenon of tidal locking. Because low-mass stars, particularly M-dwarfs, emit significantly less luminous energy than G-type stars like our Sun, their respective habitable zones are located much closer to the star.[23] Planets orbiting within these proximate habitable zones are subjected to intense gravitational gradients across their diameters. Over millions of years, tidal dissipation forces the planet's rotational period to synchronize with its orbital period.[24] Consequently, the planet presents a single hemisphere to the star in perpetual daylight, while the opposite hemisphere remains in eternal darkness.[4]

The mathematical framework governing the evolution of a planet's spin state toward synchronous rotation is highly complex. The spin state is tracked using the substellar longitude, denoted as γ, which represents the angle between the planet's long axis and a stationary inertial line relative to the mean anomaly.[23] The equation of motion for a planet's spin under the influence of tidal forces incorporates an eccentricity-dependent factor and a triaxiality factor of the principal moments of inertia. To determine the rate at which the spin state slows down to a locked state, models incorporate constant time-lag tidal dissipation, governed by the differential equation:

$$\ddot{\gamma} + \frac{1}{2}\omega_s^2 \sin(2\gamma) + \dot{n} + \epsilon\dot{\gamma} = 0$$

where the constant ϵ consolidates several physical parameters including the mass of the star, the mass of the planet, the planetary radius, the orbital semi-major axis, and the tidal parameter Q.[23] The tidal parameter varies drastically depending on the planet's composition; a Q value of 10 is typically applied to planets with thick atmospheres or global oceans, whereas a value of 100 is applied to worlds covered in ice sheets.[23]

Alternatively, the timescale required for a planet to achieve tidal locking can be estimated using a simplified physical approximation:

$$t_{lock} \approx 1486 \left(a^{0.027} M^{1/3}\right)^6$$

where t_{lock} is measured in seconds, M represents the stellar mass in grams, and a is the orbital distance in centimeters.[26] A digital twin focusing on M-dwarf systems must integrate these temporal mechanics to ascertain whether a planet has already reached a synchronous state, or if it remains in a transient, slowly rotating regime.[24] Understanding this state is paramount because a slowly rotating planet experiences Coriolis forces that fundamentally alter global atmospheric circulation, whereas a tidally locked planet relies entirely on day-to-night substellar divergence to transport heat.[28]

Atmospheric Circulation and Ocean Heat Transport

The assumption that the dayside of a tidally locked planet would boil while the nightside freezes ignores the profound impact of atmospheric and oceanic fluid dynamics. In reality, a dense atmosphere or a global ocean will transport thermal energy from the irradiated dayside to the dark nightside, radically altering the planet's surface temperature distribution.[4]

Climate feedbacks play a dictatorial role in this energy distribution. The ice-albedo feedback is a prominent positive feedback loop where cooling leads to increased ice formation; since ice possesses a high albedo, it reflects more starlight, leading to further cooling and potential global glaciation.[4] Conversely, the carbonate-silicate cycle acts as a negative, stabilizing feedback over geological timescales by sequestering or outgassing carbon dioxide in response to temperature fluctuations.[4]

When global oceans are introduced into a planetary model, the transport of heat becomes highly localized. High-resolution 3D GCM studies have demonstrated that tidally locked aquaplanets exhibit distinct climatic states based on the efficiency of Ocean Heat Transport (OHT).[25] A planet with poor OHT will enter an "eyeball" state, characterized by a circular region of liquid water precisely at the substellar point, surrounded by a global crust of thick sea ice.[25] However, when OHT is highly efficient, equatorial currents transport heat eastward, melting the ice and resulting in a "lobster" state—a configuration featuring a vast, elongated equatorial ocean that significantly expands the habitable surface area.[25] Simulating these phase transitions and thermodynamic flows requires solving the Navier-Stokes equations and thermodynamic energy transport equations across a spherical grid, an endeavor perfectly suited for acceleration via deep learning.[10]

Machine Learning Architectures in Exoclimatology

The application of machine learning to planetary science is not merely an exercise in data fitting; it requires the careful selection of architectures capable of respecting physical laws while interpolating across high-dimensional parameter spaces. For data science practitioners

participating in rapid development environments like hackathons, selecting the correct algorithm dictates whether a digital twin converges into a functional prototype or fails due to computational exhaustion.[11] The following architectures represent the cutting edge of algorithmic research in exoplanetary emulation.

Extreme Learning Machines for Surrogate Modeling

Classical General Circulation Models (GCMs), such as the ROCKE-3D framework developed by the NASA Goddard Institute for Space Studies, resolve the physical equations of atmospheric processes across highly detailed spatial and temporal grids.[2] Because these models are prohibitively slow for the statistical ensemble studies required by modern exoplanet observation, researchers have developed surrogate models to emulate GCM outputs instantaneously.[11]

Extreme Learning Machines (ELMs) have emerged as highly efficient non-gradient-based alternatives to deep neural networks for this purpose. Unlike traditional multilayer perceptrons that rely on backpropagation and gradient descent to update weights, ELMs utilize a single layer of hidden nodes where the input weights and biases are randomly assigned and remain fixed.[11] The output weights are then determined analytically in a single step via the Moore-Penrose generalized inverse.[11]

This architecture provides a drastic reduction in computational training overhead. In experiments emulating short-wave two-stream atmospheric radiative transfer—a low-dimensional sequential computation—ELMs demonstrated a training time 100,000 times faster than benchmark Bi-Directional Recurrent Neural Networks (BIRNNs) and a prediction speed 40 times faster.[11] Furthermore, ELMs proved highly sample-efficient, achieving comparable accuracy while utilizing less than one percent of the training data required by gradient-based models.[11]

For high-dimensional spatial processing, such as mapping a simulated stellar surface to a light-curve conversion, a single ELM is mathematically insufficient to capture complex spatial relationships.[11] However, an ensemble architecture comprising dozens of individual ELM predictors can achieve accuracy comparable to Convolutional Neural Networks (CNNs). While the inference speed of an ELM ensemble may lag slightly behind a highly optimized CNN, the elimination of hyperparameter tuning and iterative training renders the ELM ensemble a highly strategic choice for rapid prototyping during time-constrained computational sprints.[11]

Physics-Informed Neural Networks (PINNs)

Purely data-driven machine learning models are fundamentally agnostic to the physical laws governing the universe. Consequently, when asked to predict planetary temperatures or atmospheric transport outside the specific distribution of their training data, they often produce physically impossible results.[33] Physics-Informed Neural Networks (PINNs) solve this

exact problem by embedding partial differential equations (PDEs) directly into the neural network's loss function.[33]

In a standard neural network, the objective is to minimize the Mean Squared Error (MSE) between the predicted outputs and the ground truth data. In a PINN, the loss function is expanded to penalize the network if its predictions violate the governing laws of physics.[33] The total loss function of a PINN is formally constructed as:

$$L_{total} = \lambda_{data} L_{data} + \lambda_{PDE} L_{PDE} + \lambda_{BC} L_{BC} + \lambda_{IC} L_{IC}$$

where L_{data} represents the error relative to known observational data, L_{PDE} computes the residual of the governing differential equations (such as the heat transfer equation or the Navier-Stokes equations) using automatic differentiation, L_{BC} enforces spatial boundary conditions, and L_{IC} enforces initial temporal conditions.[36] The coefficients denoted by λ dynamically weight the importance of each constraint during optimization.

PINNs have been successfully deployed to model radiative transfer (RT) processes in exoplanetary atmospheres.[35] Traditional fast RT models often simplify light scattering phenomena by treating scattering entirely as absorption, resulting in severe inaccuracies when analyzing clouds or hazes.[35] By structuring a PINN to compute the exact RT differential equations—including pressure-dependent coefficients for Rayleigh scattering—researchers can predict the transmission spectra of exoplanets addressing both direct and diffuse stellar light.[35] Furthermore, PINNs are incredibly effective at solving the heat diffusion equation across a simulated planetary surface, making them ideal for modeling the temperature distribution and subsurface ocean formation on tidally locked worlds like Proxima Centauri b.[39]

Bayesian Frameworks and Uncertainty Quantification

Because observational data for exoplanets is sparse and subject to high instrument noise, deterministic machine learning models can exhibit unwarranted confidence in erroneous predictions. Bayesian frameworks address this by treating model parameters and outputs as probability distributions rather than point estimates, inherently quantifying the uncertainty of a prediction.[2]

A sophisticated application of this paradigm is the use of Bayesian Adaptive Exploration combined with spline curves to predict atmospheric absorption spectra.[2] In this framework, a predictive forward model is trained on a combination of real solar system spectra and synthetic data generated by the ROCKE-3D GCM. The algorithm models atmospheric spectra across discrete bins from the visible to near-infrared ranges, treating each bin height as a non-linear function of dozens of planetary parameters.[2] Once established, the forward model operates within a Bayesian inference engine to perform atmospheric retrieval—the inverse operation of

estimating the most probable planetary parameters, including day-side temperature and atmospheric composition, based strictly on an observed absorption spectrum.[2]

Crucially, Bayesian Adaptive Exploration algorithms actively analyze the 30-dimensional planetary parameter space to identify regions where model uncertainty is highest.[2] This dynamic targeting dictates where expensive 3D GCMs should run new simulations, effectively creating an active learning loop that maximizes computational efficiency while continuously refining the planetary digital twin.[2]

Handling Class Imbalance with Generative Adversarial Networks

One of the most profound obstacles in exoplanetary data science is the severe class imbalance inherent in astronomical catalogs. In the NASA Exoplanet Archive, out of thousands of confirmed exoplanets, only a minute fraction possess the physical parameters necessary to be classified as potentially habitable.[1] When standard classification algorithms, such as Logistic Regression, are trained on this skewed data, they often achieve high overall accuracy simply by predicting that every planet is non-habitable, resulting in zero sensitivity for the minority class.[1]

While techniques such as undersampling the majority class or applying Elastic Gini Criteria in Decision Trees offer marginal improvements, Generative Adversarial Networks (GANs) provide a transformative solution.[1] Specifically, Conditional Generative Adversarial Networks (cGANs) consist of two competing neural architectures: a generator that synthesizes artificial planetary data, and a discriminator that attempts to distinguish between the synthesized data and actual astronomical observations.[12]

By conditioning the generator on the specific parameters of the habitable minority class, the cGAN learns the underlying multivariate distribution of habitability. It then synthesizes highly realistic data points representing hypothetical habitable exoplanets, effectively inflating the minority class to achieve parity with the non-habitable class.[12] When custom classifiers, such as XGBoost or sophisticated Random Forests, are subsequently trained on this robust, balanced dataset, their capacity to accurately classify borderline planetary candidates improves dramatically, frequently exceeding predictive accuracies of 96 percent.[5]

Data Ecosystems and Feature Engineering

The foundation of any machine learning model is the quality, structure, and pre-processing of its ingested data. For astrobiological digital twins, researchers primarily rely on observational catalogs for classification tasks and synthetic simulation grids for emulation tasks.[7]

The NASA Exoplanet Archive and Observational Datasets

The NASA Exoplanet Archive serves as the premier repository for stellar and planetary parameters, aggregating data from transit surveys like Kepler and TESS, radial velocity

measurements, and direct imaging campaigns.[42] Accessing this data programmatically is typically achieved via the Table Access Protocol (TAP) service or the application programming interface (API), which allows researchers to query millions of light curves and thousands of confirmed planetary system parameters using Python scripting.[45]

The dataset contains nearly a hundred columns of raw astrophysical data per entry, necessitating rigorous feature engineering before deployment in a machine learning pipeline.[46] Crucial primary features include the orbital semi-major axis, the planetary radius, planetary mass, insolation flux, and the host star's effective temperature.[2] However, the most informative features are often composite metrics derived mathematically from the raw data.

The Earth Similarity Index (ESI) is the preeminent metric for gauging the physical resemblance of an exoplanet to Earth. The ESI is a multi-parameter distance metric scaling from 0 to 1, where 1 represents an identical Earth analog.[47] It is mathematically subdivided into an Interior ESI and a Surface ESI. The fundamental calculation for any single property utilizes a weighted formula:

$$ESI_i = \left(1 - \left|\frac{x_i - x_{i,0}}{x_i + x_{i,0}}\right|\right)^{w_i}$$

where x_i is the measured parameter of the exoplanet, $x_{i,0}$ is the terrestrial reference value, and w_i is the specific weight exponent for that property.[48] The Interior ESI combines the planet's mean radius and bulk density, while the Surface ESI evaluates the escape velocity and surface temperature.[48] Because the scale is highly sensitive to thermodynamic state changes, surface temperature is assigned the highest weight exponent in the calculation.[48] Algorithms that incorporate the ESI alongside raw parameters demonstrate significantly higher confidence and interpretability when scoring candidates for potential follow-up observation.[41]

Table 1 provides a mathematical breakdown of the components governing the Earth Similarity Index, demonstrating how feature engineering translates raw astronomical observations into structured inputs for machine learning classifiers.

Planetary Property	**ESI Sub-Category**	**Terrestrial Reference Value (xi,0)**	**Weight Exponent (wi)**	**Physical Implication for Habitability**
Mean Radius	Interior	1.0 Earth Radii	0.57	Determines gravity and atmospheric retention

				capacity.
Bulk Density	Interior	1.0 Earth Density	1.07	Indicates rocky composition versus gaseous envelope.
Escape Velocity	Surface	1.0 Earth Units	0.70	Dictates the stability of an atmosphere against solar wind stripping.
Surface Temperature	Surface	288.0 Kelvin	5.58	Highest weighted parameter; critical boundary for the existence of liquid water.

Table 1: Reference values and weight exponents for the calculation of the Earth Similarity Index (ESI), utilized as a composite feature in machine learning habitability pipelines.[48]

ROCKE-3D Simulation Grids and NetCDF Frameworks

While observational archives provide the parameters of discovery, General Circulation Model outputs provide the complex meteorological and thermodynamic data required to build a digital twin. The ROCKE-3D GCM datasets, hosted by the Goddard Institute for Space Studies, are paramount for this task.[7]

The ROCKE-3D framework simulates rocky planet climates across a standardized grid, typically resolving at 4 by 5 degrees in latitude and longitude, featuring numerous atmospheric layers extending up to 0.1 hPa.[29] The archived output files are structured in the Network Common Data Form (NetCDF), a set of software libraries and data formats that support the creation, access, and sharing of array-oriented scientific data.[52]

For machine learning applications targeting tidally locked systems, researchers can extract specific continuous simulation variables representing the planet's state.[7] These multi-dimensional arrays include boundary topography, stellar heating (W/m^2), specific

humidity, sea ice fraction, cloud cover percentages, and planetary albedo.[7] By parsing these vast NetCDF files, machine learning engineers can structure high-fidelity tensor datasets where the input features are the underlying planetary assumptions (e.g., synchronous CO2 atmosphere, pre-industrial Earth, anoxic environments) and the targets are the resulting spatially resolved climate variables.[7]

To properly view the detailed structural data available through the NASA systems, Table 2 delineates a selection of specific features that machine learning models must map to understand climate variance.

Variable Category	Continuous Features Extracted from GCM Grids	Utility in Machine Learning Emulation
Thermodynamic	Surface Temperature, Absorbed Starlight, Emitted Heat	Predicting the boundaries of the habitable terminator line.[7]
Hydrological	Precipitation rate, Column Water Vapor, Soil Moisture	Tracking state changes and modeling the "lobster" ocean current state.[7]
Atmospheric	Sea Level Pressure, 500 mb Height, Cloud Cover (%)	Assessing the viability of the runaway greenhouse limit.[7]
Cryospheric	Sea Ice Fraction, Snow Depth	Evaluating the strength of the positive ice-albedo feedback loop.[7]

Table 2: Key physical features mapped within the ROCKE-3D GCM NetCDF outputs utilized for training digital twin surrogate models.[7]

Strategic Project Blueprints for Computational Astrobiology

The theoretical mechanics, advanced algorithmic architectures, and rich datasets discussed above provide the foundation for executing highly complex, innovative projects. For students with strong mathematical and programming capabilities—such as those at the WUT Faculty of Mathematics and Information Science—the challenge is applying these advanced concepts within a severely constrained timeframe. The following blueprints are designed specifically for

Category 3 of the HACK-4-SAGES initiative ("Digital Twins in Exoplanet Habitability").[4] They align with the 72-hour development window, rely entirely on accessible open-source data, and leverage cutting-edge AI methodologies to advance the state of astrobiological modeling.[53]

Blueprint Alpha: PINN-Based Spatiotemporal Heat Distribution Model for Tidally Locked Worlds

The behavior of temperature and water distribution on a tidally locked, synchronous planet represents a pinnacle problem in exoclimatology.[25] Traditional 1D energy balance models simulate temperature conceptually, but fail to account for continuous spatial gradients and longitudinal variations.[55] Conversely, running a full 3D GCM to test small variations in a planet's atmospheric composition is too slow for real-time digital twin interaction.[10]

This project proposes the construction of a Physics-Informed Neural Network (PINN) designed to solve the 2D spherical heat diffusion equation across the surface of a tidally locked rocky exoplanet. The objective is to rapidly emulate the temperature gradients spanning from the substellar point (eternal day) to the antistellar point (eternal night), accurately identifying the width and stability of the habitable terminator line.[4]

The architecture utilizes a multi-layer perceptron where the inputs are spatial coordinates (latitude and longitude) and time, and the output is the localized surface temperature. Crucially, the loss function is mathematically bound to the physical laws of thermal diffusion. The algorithm applies automatic differentiation to penalize the network if its predicted temperature gradients violate the heat equation, ensuring that the emulation adheres to the principles of energy conservation.[33]

The boundary conditions enforced within the PINN will integrate the incident stellar flux from an M-dwarf star, simulating the perpetual irradiation of the dayside.[30] Furthermore, the model can incorporate a parameterized constant representing wind-driven heat redistribution, simulating how an atmosphere shifts the hottest point of the planet slightly eastward of the substellar point.[10]

By compiling this model into a digital twin interface using a Python dashboard framework like Streamlit, users could dynamically manipulate the planetary albedo or the stellar distance via a slider.[46] The PINN would instantly converge on the new steady-state temperature distribution, bypassing the need for computationally heavy differential equation solvers. The success of this model is evaluated by its ability to accurately identify surface regions maintaining temperatures between 0°C and 100°C, providing a real-time habitability fraction map.[4] This project requires profound knowledge of partial differential equations and PyTorch optimization, mapping perfectly to the curriculum of advanced applied mathematics students.

Blueprint Beta: Extreme Learning Machine (ELM) Ensemble for Rapid

Radiative Transfer Emulation

A core challenge in assessing habitability is determining how an exoplanet's atmosphere filters and traps heat through the greenhouse effect. Calculating the exact radiative transfer of energy through atmospheric layers is a severe computational bottleneck in traditional simulations, often requiring complex iterative integrals over the entire electromagnetic spectrum.[11]

This project entails building a surrogate model using an ensemble of Extreme Learning Machines (ELMs) to bypass traditional radiative transfer calculations. The goal is to ingest high-dimensional tabular data from the ROCKE-3D GCM and instantly output the atmospheric vertical temperature profiles and simulated transmission spectra.[2]

The methodology requires parsing the ROCKE-3D NetCDF datasets, specifically isolating the atmospheric compositions of simulated worlds, such as Archean Earth environments or hypothetical synchronous CO2 atmospheres.[7] Because the input data involves complex vertical column approximations, a single neural network is prone to overfitting and spatial degradation. Therefore, an ensemble of ELMs is initialized.[11] Because ELMs assign hidden layer weights randomly and solve for output weights analytically rather than through epochs of gradient descent, the entire ensemble can be trained on a standard local machine in a matter of seconds.[11] This architectural advantage is perfectly suited for a three-day hackathon sprint where time spent training models is a critical liability.

The ensemble's output will emulate the discrete bins of atmospheric absorption spectra, mirroring the expected observations from instruments aboard the JWST, such as the Near-Infrared Spectrograph (NIRSpec) and the Mid-Infrared Instrument (MIRI).[2] By bridging the gap between theoretical atmospheric composition and expected observational spectra, this ELM ensemble acts as a digital twin mechanism that allows astronomers to quickly hypothesize what atmospheric signatures might correspond to a habitable climate.

Blueprint Gamma: cGAN-Augmented Bayesian Framework for Probabilistic Habitability Classification

The vast majority of planets identified in the NASA Exoplanet Archive are uninhabitable gas giants or sterile, irradiated rocks. Building a classification model to identify promising candidates for follow-up observation is notoriously difficult when the target class (habitable worlds) is statistically negligible.[1]

This blueprint proposes a hybrid data engineering and classification pipeline. The initial phase utilizes a Conditional Generative Adversarial Network (cGAN) to synthesize high-fidelity artificial data representing habitable Earth-like planets. By training the cGAN on the statistical boundaries of known temperate rocky worlds—leveraging constraints like the equilibrium temperature and rocky density thresholds—the generator learns to output realistic

combinations of planetary radii, equilibrium temperatures, and stellar insolation fluxes.[12] The dataset is thus artificially balanced, removing the bias that typically cripples standard logistic regression models by causing them to blindly predict the majority class.[1]

Following dataset augmentation, engineered features—specifically the global Earth Similarity Index (ESI) and the Thermal Suitability Score (TSS)—are calculated for both the real and synthetic planets using Python libraries like Pandas and Scikit-learn.[1] The classification engine is then constructed using an optimized XGBoost algorithm, a decision-tree-based framework highly effective at capturing non-linear feature interactions.[5]

To elevate this project beyond a simple binary classifier, it incorporates Bayesian probabilities. Instead of outputting a rigid "habitable" or "non-habitable" label, the XGBoost classifier is calibrated to output a probability distribution, allowing the digital twin to rank planets by the statistical likelihood of habitability.[57] This aligns with the epistemological reality of astrobiology: habitability is a spectrum, not a binary state.[59] The resulting digital twin acts as a rigorous filtering mechanism, allowing researchers to input the raw data of a newly discovered transit event via the NASA API and instantly receive a probabilistic assessment of its habitability, drastically reducing the time required for manual vetting.[22]

Blueprint Delta: Deep Emulator for Ocean Heat Transport and Ice-Albedo Feedbacks

The presence of surface water fundamentally changes the thermodynamic behavior of a planet. Tidally locked planets exhibit a delicate balance between atmospheric collapse on the nightside and runway greenhouse effects on the dayside, a balance largely mediated by Ocean Heat Transport (OHT).[25]

This project aims to build a deep learning emulator that maps the intricate fluid dynamics of water phase changes—specifically predicting whether a tidally locked aquaplanet will stabilize into an isolated "eyeball" state or a widely circulated "lobster" state.[25] The model relies on a dense neural network (DNN) trained to emulate the outputs of a coupled atmosphere-ocean general circulation model, utilizing specific sub-grids derived from the ROCKE-3D framework.[10]

The input parameters involve planetary rotation rate (simulating the transition from a 3:2 orbital resonance to a fully synchronous state), global water inventory, and incident stellar flux.[4] The network is tasked with predicting the fractional coverage of sea ice and the specific latitudinal locations of open oceanic water.[4]

To capture the non-linear tipping points caused by the ice-albedo positive feedback loop, the neural network incorporates specialized activation functions and recurrent time-stepping mechanisms. This allows the digital twin to simulate temporal evolution. A user interacting with this digital twin could gradually decrease the simulated star's luminosity, watching in real-time as the neural network predicts the advancing ice sheets, mapping the exact thermodynamic

tipping point where the planet undergoes rapid, irreversible global glaciation.[4]

Table 3 synthesizes these proposed architectural blueprints, outlining their specific computational methodologies and their direct alignment with the core challenges of exoplanet characterization.

Blueprint Designation	Core ML Architecture	Primary Astrobiological Objective	Computational Advantage for Hackathons
Alpha: Heat Distribution	Physics-Informed Neural Network (PINN)	Modeling 2D spherical heat diffusion and terminator habitability.	Bypasses numerical PDE solvers, ensuring rapid, physics-compliant emulation.[37]
Beta: Radiative Transfer	Extreme Learning Machine (ELM) Ensemble	Emulating vertical atmospheric temperature profiles and absorption spectra.	Eliminates iterative gradient-descent training, allowing instantaneous deployment.[11]
Gamma: Habitability Classification	cGAN augmented XGBoost Pipeline	Classifying exoplanet viability via Earth Similarity Index (ESI) matrices.	Solves severe dataset class imbalance; outputs probabilistic certainty.[12]
Delta: Ocean Transport	Deep Neural Network (DNN)	Predicting "eyeball" vs "lobster" ocean states and ice-albedo tipping points.	Emulates complex GCM fluid dynamics in real-time, enabling interactive temporal simulations.[10]

Table 3: Strategic overview of advanced machine learning blueprints for digital twins in exoplanetary habitability.

Methodological Considerations and Performance Evaluation

The deployment of machine learning in the physical sciences demands rigorous validation protocols. Because these models aim to approximate reality in environments where direct empirical verification is currently impossible (i.e., we cannot physically launch a probe to Proxima Centauri b to verify a surface temperature map), their internal mathematical consistency is paramount. For hackathon participants, proving the mathematical validity of a model is often more heavily weighted by judges than deploying a visually appealing but scientifically inaccurate dashboard.

When developing surrogate models, particularly Deep Neural Networks and Extreme Learning Machines meant to emulate 3D GCMs, the primary evaluation metric is the Mean Absolute Error (MAE) and the Root Mean Squared Error (RMSE) comparing the emulation to the baseline ROCKE-3D output.[61] Emulators must demonstrate that their predicted gas temperatures remain within highly specific tolerances—often within tens of parts per million—to ensure that secondary chemical predictions, such as cloud formation, are not invalidated by underlying temperature artifacts.[10]

For Physics-Informed Neural Networks, evaluation extends beyond point-wise accuracy against a dataset. The model must be subjected to error analysis within the space of smooth functions, demonstrating that the difference between the exact analytical solution of the differential equation and the predicted solution is strictly bounded by the physics-informed loss functional.[40] If the PINN's residual loss is minimized successfully, researchers can trust the model's extrapolation capabilities far more than a purely data-driven network, as the PINN cannot easily predict a state that violates the conservation of energy.[62]

For classification pipelines handling severe class imbalance, such as determining habitability from the Kepler dataset, overall accuracy is a highly deceptive metric. A model predicting a 99% accuracy might achieve this simply by failing to identify the 1% of habitable planets in the dataset, effectively acting as a null classifier.[1] Therefore, evaluations must rely on the Macro F1 score, which calculates the harmonic mean of precision and recall for each class independently, ensuring that the model maintains high sensitivity for the critical minority class of temperate, rocky worlds.[1]

Finally, the computational architecture itself must be considered in the context of event constraints. Advanced hackathon projects must optimize their training routines to run on accessible hardware within a 72-hour window.[54] The use of ELMs, parameterized PINNs, and efficient tree-based algorithms (XGBoost) specifically mitigates the need for massive cloud-computing infrastructure or weeks of GPU training time. This algorithmic efficiency allows complex astrobiological digital twins to be hosted and operated smoothly on standard consumer hardware or interactive web interfaces, fulfilling the fundamental requirement of a responsive digital twin.[11]

Conclusion

The intersection of astrophysics, applied mathematics, and machine learning represents the absolute frontier of habitability research. As observational missions continue to inundate the scientific community with photometric and spectroscopic data, the bottleneck has irreversibly shifted from data acquisition to data interpretation. Digital twins are the necessary evolution of planetary modeling, providing the interactive, high-fidelity environments required to simulate complex climatic feedbacks, atmospheric radiative transfer, and oceanic fluid dynamics.

By replacing computationally prohibitive General Circulation Models with advanced machine learning architectures—ranging from Physics-Informed Neural Networks that respect thermodynamic laws to Generative Adversarial Networks that resolve statistical imbalances—researchers can simulate millions of years of planetary evolution in milliseconds. For computational initiatives such as Category 3 of HACK-4-SAGES, these data-driven frameworks offer a pathway to execute theoretically rigorous, highly innovative projects within tightly constrained timeframes. Ultimately, the synthesis of machine learning and exoclimatology does not merely accelerate computation; it fundamentally enhances the analytical capacity of astrobiologists to identify the elusive signatures of life across the cosmos.

Cytowane prace

1. (PDF) PREDICTING THE HABITABILITY OF EXOPLANETS USING ..., otwierano: lutego 18, 2026, https://www.researchgate.net/publication/394583617_PREDICTING_THE_HABITABILITY_OF_EXOPLANETS_USING_STATISTICAL_MODELING_AND_MACHINE_LEARNING
2. Understanding Exoplanet Habitability: A Bayesian ML Framework for Predicting Atmospheric Absorption Spectra - MDPI, otwierano: lutego 18, 2026, https://www.mdpi.com/2673-9984/12/1/9
3. Trehan et al. 2025: Understanding exoplanet habitability: A Bayesian ML framework for predicting atmospheric... - NASA GISS, otwierano: lutego 18, 2026, https://www.giss.nasa.gov/pubs/abs/tr09300n.html
4. Hackathon_Questions.pdf
5. Habitable exoplanet - a statistical search for life - Frontiers, otwierano: lutego 18, 2026, https://www.frontiersin.org/journals/astronomy-and-space-sciences/articles/10.3389/fspas.2025.1674754/full
6. Digital Twins of the Earth System via Hybrid Physics-AI Models - SIAM.org, otwierano: lutego 18, 2026, https://www.siam.org/publications/siam-news/articles/digital-twins-of-the-earth-system-via-hybrid-physics-ai-models/
7. Data.GISS: Simulations of Planetary Climates with ROCKE-3D - NASA, otwierano: lutego 18, 2026, https://data.giss.nasa.gov/rocke3d/maps/
8. Enhancing 3D Planetary Atmosphere Simulations with a Surrogate Radiative Transfer Model - arXiv, otwierano: lutego 18, 2026, https://arxiv.org/html/2407.08556v1
9. ROCKE-3D 2.0: an updated general circulation model for simulating the climates

of rocky planets - GMD, otwierano: lutego 18, 2026, https://gmd.copernicus.org/articles/18/5825/2025/
10. [2508.10827] Accelerating exoplanet climate modelling: A machine learning approach to complement 3D GCM grid simulations - arXiv, otwierano: lutego 18, 2026, https://arxiv.org/abs/2508.10827
11. Extreme learning machines for exoplanet simulations: a faster ..., otwierano: lutego 18, 2026, https://academic.oup.com/rasti/article/doi/10.1093/rasti/rzaf050/8305079
12. Predicting Exoplanet Habitability Using Machine Learning - NORMA@NCI Library, otwierano: lutego 18, 2026, https://norma.ncirl.ie/8541/
13. HACK-4-SAGES – Digital Twins in Astrobiology Hackathon, otwierano: lutego 18, 2026, https://hack-4-sages.org/
14. HACK 4 SAGES: Digital Twins in Astrobiology – call for students and mentors - DEADLINE 23.02.2026 / Important information / News / Home - CWM PW, otwierano: lutego 18, 2026, https://www.cwm.pw.edu.pl/en/News/Important-information/HACK-4-SAGES-Digital-Twins-in-Astrobiology-call-for-students-and-mentors-DEADLINE-23.02.2026
15. Planet Habitability of Exoplanets - NYU Abu Dhabi, otwierano: lutego 18, 2026, https://nyuad.nyu.edu/en/events/2021/march/planet-habitability-of-exoplanets.html
16. How Artificial Intelligence Will Help Humans to Study The Habitable Planet in The Universe - ijsred.com, otwierano: lutego 18, 2026, https://ijsred.com/volume8/issue5/IJSRED-V8I5P165.pdf
17. Eccentric orbits may enhance the habitability of Earth-like exoplanets - Oxford Academic, otwierano: lutego 18, 2026, https://academic.oup.com/mnras/article/532/4/4511/7717378
18. Habitability and sub glacial liquid water on planets of M-dwarf stars - PMC - NIH, otwierano: lutego 18, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC10104869/
19. Habitability of terrestrial-mass planets in the HZ of M Dwarfs – I. H/He-dominated atmospheres | Monthly Notices of the Royal Astronomical Society | Oxford Academic, otwierano: lutego 18, 2026, https://academic.oup.com/mnras/article/459/4/4088/2624070
20. HABITABLE ZONES AROUND MAIN-SEQUENCE STARS: NEW ESTIMATES (vol 765, pg 131, 2013) - ResearchGate, otwierano: lutego 18, 2026, https://www.researchgate.net/publication/258259257_HABITABLE_ZONES_AROUND_MAIN-SEQUENCE_STARS_NEW_ESTIMATES_vol_765_pg_131_2013
21. [1705.10362] Habitable Moist Atmospheres On Terrestrial Planets Near the Inner Edge Of the Habitable Zone Around M-dwarfs - arXiv.org, otwierano: lutego 18, 2026, https://arxiv.org/abs/1705.10362
22. Machine Learning For Exoplanet Discovery: Validating TESS Candidates and Identifying Planets in the Habitable Zone - Astrobiology, otwierano: lutego 18, 2026, https://astrobiology.com/2025/12/machine-learning-for-exoplanet-discovery-validating-tess-candidates-and-identifying-planets-in-the-habitable-zone.html
23. Day and Night: Habitability of Tidally Locked Planets with Sporadic ..., otwierano:

lutego 18, 2026, https://arxiv.org/abs/2303.14546
24. [1708.02981] Tidal Locking of Habitable Exoplanets - arXiv, otwierano: lutego 18, 2026, https://arxiv.org/abs/1708.02981
25. Role of ocean heat transport in climates of tidally locked exoplanets around M dwarf stars, otwierano: lutego 18, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC3896211/
26. Tidal Lock Radius in Habitable Zones - Physics Stack Exchange, otwierano: lutego 18, 2026, https://physics.stackexchange.com/questions/12541/tidal-lock-radius-in-habitable-zones
27. Day 'N' Nite: Habitability of Tidally Locked Planets with Sporadic Rotation - OAsis Research Repository - University of Nevada, Las Vegas, otwierano: lutego 18, 2026, https://oasis.library.unlv.edu/cgi/viewcontent.cgi?article=5622&context=thesesdissertations
28. Day and night: habitability of tidally locked planets with sporadic rotation - Oxford Academic, otwierano: lutego 18, 2026, https://academic.oup.com/mnras/article/524/4/5708/7230370
29. Climates of Warm Earth-like Planets I: 3-D Model Simulations ..., otwierano: lutego 18, 2026, https://ntrs.nasa.gov/api/citations/20180007568/downloads/20180007568.pdf
30. A One-Dimensional Energy Balance Model Parameterization for the Formation of CO2 Ice on the Surfaces of Eccentric Extrasolar Planets - Astrobiology, otwierano: lutego 18, 2026, https://astrobiology.com/2025/01/a-one-dimensional-energy-balance-model-parameterization-for-the-formation-of-co2-ice-on-the-surfaces-of-eccentric-extrasolar-planets.html
31. ROCKE-3D 2.0: An updated general circulation model for simulating the climates of rocky planets - EGUsphere, otwierano: lutego 18, 2026, https://egusphere.copernicus.org/preprints/2025/egusphere-2025-925/
32. Surrogate model for geological CO2 storage and its use in hierarchical MCMC history matching, otwierano: lutego 18, 2026, https://sccs.stanford.edu/sites/g/files/sbiybj17761/files/media/file/surrogate-model-for-geological-co2-storage-and-its-use-in-hierarchical-mcmc-history-matching.pdf
33. Solving the Heat Equation with Physics-Informed Neural Networks: A Practical Guide | by Nikolaos Pallikarakis, PhD | Feb, 2026 | Medium, otwierano: lutego 18, 2026, https://medium.com/@pallikarakis.n/solving-the-heat-equation-with-physics-informed-neural-networks-a-practical-guide-25df376f1b04
34. Approximating Rayleigh Scattering in Exoplanetary Atmospheres using Physics-informed Neural Networks (PINNs) - ResearchGate, otwierano: lutego 18, 2026, https://www.researchgate.net/publication/382851373_Approximating_Rayleigh_Scattering_in_Exoplanetary_Atmospheres_using_Physics-informed_Neural_Network

s_PINNs
35. [2408.00084] Approximating Rayleigh Scattering in Exoplanetary Atmospheres using Physics-informed Neural Networks (PINNs) - arXiv, otwierano: lutego 18, 2026, https://arxiv.org/abs/2408.00084
36. Physics-Informed Neural Networks for the Heat Equation with Source Term under Various Boundary Conditions - MDPI, otwierano: lutego 18, 2026, https://www.mdpi.com/1999-4893/16/9/428
37. nkusla/heat-transfer-pinn: Simulating heat propagation using Physics-informed Neural Network - GitHub, otwierano: lutego 18, 2026, https://github.com/nkusla/heat-transfer-pinn
38. Approximating Rayleigh Scattering in Exoplanetary Atmospheres using Physics-informed Neural Networks (PINNs) - Astrobiology, otwierano: lutego 18, 2026, https://astrobiology.com/2024/08/approximating-rayleigh-scattering-in-exoplanetary-atmospheres-using-physics-informed-neural-networks-pinns.html
39. Simulation of Heat Diffusion in a Tidally-Locked Exoplanet - AWS, otwierano: lutego 18, 2026, https://terra-docs.s3.us-east-2.amazonaws.com/IJHSR/Articles/volume4-issue3/2022_43_p17_Murgia.pdf
40. Physics-informed neural network for the heat equation under imperfect contact conditions and its error analysis - AIMS Press, otwierano: lutego 18, 2026, https://www.aimspress.com/article/doi/10.3934/math.2025364?viewType=HTML
41. Classifying Exoplanet Habitability using Machine Learning | by Joshua Araujo - Medium, otwierano: lutego 18, 2026, https://joshua-araujo.medium.com/classifying-exoplanet-habitability-using-machine-learning-d1497d9d7547
42. NASA Exoplanet Archive, otwierano: lutego 18, 2026, https://exoplanetarchive.ipac.caltech.edu/
43. [OC] I Visualized the Orbits of 4,000+ Real Exoplanets from NASA's Archives in an Interactive 3D Universe : r/dataisbeautiful - Reddit, otwierano: lutego 18, 2026, https://www.reddit.com/r/dataisbeautiful/comments/1npmod2/oc_i_visualized_the_orbits_of_4000_real/
44. NASA Exoplanet Archive Overview and Holdings - Caltech, otwierano: lutego 18, 2026, https://exoplanetarchive.ipac.caltech.edu/docs/intro.html
45. Using the Application Programming Interface (API) - NASA Exoplanet Archive - Caltech, otwierano: lutego 18, 2026, https://exoplanetarchive.ipac.caltech.edu/docs/program_interfaces.html
46. ShubhamP1028/Exoplanet-Habitability-Analysis-and-Prediction-Model: This project performs a comprehensive analysis of exoplanet datasets using Python. It involves preprocessing NASA's Kepler/K2 mission data, handling missing values, encoding categorical variables, and performing feature selection. EDA reveals correlations between planetary and stellar parameters Random Forest model to - GitHub, otwierano: lutego 18, 2026, https://github.com/ShubhamP1028/Exoplanet-Habitability-Analysis-and-Prediction-Model

47. What is the Earth Similarity Index (ESI) for exoplanets? - YouTube, otwierano: lutego 18, 2026, https://www.youtube.com/watch?v=VU39UWkFHqQ
48. Earth Similarity Index (ESI) - PHL @ UPR Arecibo, otwierano: lutego 18, 2026, https://phl.upr.edu/projects/earth-similarity-index-esi
49. Earth Similarity Index - Dodona, otwierano: lutego 18, 2026, https://dodona.be/en/activities/1726580002/
50. Indexing Exoplanets with Physical Conditions Potentially Suitable for Rock-Dependent Extremophiles - PMC, otwierano: lutego 18, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC7175269/
51. ROCKE-3D 2.0: An updated general circulation model for simulating the climates of rocky planets - EGUsphere, otwierano: lutego 18, 2026, https://egusphere.copernicus.org/preprints/2025/egusphere-2025-925/egusphere-2025-925-ATC1.pdf
52. simulation - Open-source earth model for climate change, otwierano: lutego 18, 2026, https://earthscience.stackexchange.com/questions/19235/open-source-earth-model-for-climate-change
53. FAQ - HACK-4-SAGES, otwierano: lutego 18, 2026, https://hack-4-sages.org/faq/
54. How to join? - HACK-4-SAGES, otwierano: lutego 18, 2026, https://hack-4-sages.org/how-to-join/
55. SWAMPE: A Shallow-Water Atmospheric Model in Python for Exoplanets - Open Journals, otwierano: lutego 18, 2026, https://www.theoj.org/joss-papers/joss.04872/10.21105.joss.04872.pdf
56. BACHELOR THESIS Atmospheric transmission simulations of exoplanet TRAPPIST-1e - DiVA, otwierano: lutego 18, 2026, http://www.diva-portal.org/smash/get/diva2:1884289/FULLTEXT01.pdf
57. Exoplanet Habitability Detector - | International Journal of Science and Research Archive, otwierano: lutego 18, 2026, https://journalijsra.com/sites/default/files/fulltext_pdf/IJSRA-2025-3335.pdf
58. Examining Exoplanets for Habitability - Jeffrey Robbins, otwierano: lutego 18, 2026, https://jtrobbins.com/portfolio/jrobbins_cmsc320.html
59. Climates of Warm Earth-like Planets III: Fractional Habitability from a Water Cycle Perspective, otwierano: lutego 18, 2026, https://ntrs.nasa.gov/api/citations/20205001008/downloads/Climates_of_Warm_Earth_like_Planets_III-accepted.pdf
60. MODELLING STUDIES OF EXOPLANETARY OCEANS - The University of East Anglia, otwierano: lutego 18, 2026, https://ueaeprints.uea.ac.uk/67833/1/2018CullumJPhD.pdf
61. A Digital Twin of a Water Distribution System by Using Graph Convolutional Networks for Pump Speed-Based State Estimation - MDPI, otwierano: lutego 18, 2026, https://www.mdpi.com/2073-4441/14/4/514
62. Physics-informed neural network for the heat equation under imperfect contact conditions and its error analysis - AIMS Press, otwierano: lutego 18, 2026, https://www.aimspress.com/article/id/67f3b285ba35de0692a8c879